

Of course — let me walk through the whole notebook in plain English.

What Part 5 is trying to answer:

"If we run SPRING20 again next year, how many orders and how much revenue should Finance budget for?"

The core problem we're solving:

We can't just look at this year's actuals and say "it'll happen again." We need to account for the fact that the business is growing, and that we're not 100% sure how much. So instead of one number, we produce a range (low / mid / high) with explicit reasons for each.

Section by section:

Sections 1–2 — Confirmed inputs and LY uplift We lock in all the numbers from your data — TY actuals, baseline averages, and the derived LY figures ($TY \times 0.75$ as the assignment instructs). We also compute the "uplift factor" — how much the promo amplified daily orders above baseline last year (~2.5x).

Section 3 — AOV adjustment The 20% discount means each order is worth less revenue than a normal day. But your data shows customers buy slightly more units during the promo (UPT went from 2.05 → 2.11), which partially offsets the discount. We model this compression at 80–92% of baseline AOV depending on scenario.

Section 4 — Scenario parameters Three scenarios defined by two dials:

- How fast is DTC growing year over year? (1.10 / 1.18 / 1.25)
- Will the promo drive the same relative lift as last year? (85% / 100% / 110%)

Section 5 — The actual forecast (two-anchor blend) This is the heart of the model. Instead of one calculation, we run two independent ones and blend them 60/40:

- Anchor A asks: *"what did LY promo do, scaled up for growth?"*
- Anchor B asks: *"what does TY baseline look like, amplified by the LY uplift factor?"*

Blending them means if last year had an unusual spike or dip, it doesn't fully distort the forecast.

Section 6 – Summary table Puts all three scenarios side by side with TY actuals for comparison – orders, revenue, AOV, daily averages.

Section 7 – Model accuracy check Since we already know what TY actuals were, we can check how wrong or right each scenario would have been. This tells Finance how much to trust the model.

Section 8 – Sensitivity analysis Asks: if we're slightly wrong about one assumption, how much does the final revenue number move? The assumption with the biggest swing is the one Finance should pressure-test hardest before committing to inventory or staffing.

Section 9 – Scenario drivers For each scenario, four business conditions that would need to be true for it to happen. Turns abstract numbers into concrete planning logic.

Section 10 – Finance brief Five bullet points auto-generated from your actual numbers – ready to paste into an email to the Finance team.

Does that make sense? Any section you want me to go deeper on before we move to Part 6?